

Class begins on Michigan Time. Before class begins, please do the following:

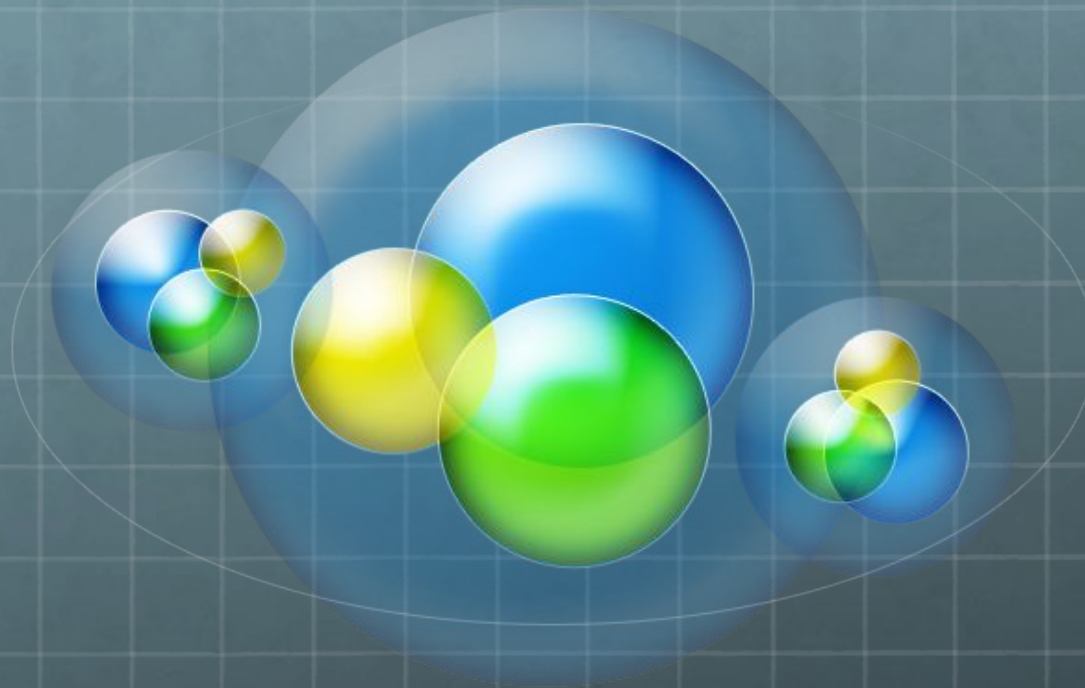
1. Visit **www.umich.edu/~cja/** and click "STATS 531 Introduction to Flux" to view these slides.
2. Login to Flux using the instructions on the first few slides.



STATS 531

Introduction to Flux

Dr Charles J Antonelli, LSA IT ARS
Mark Montague, LSA IT ARS
February, 2018



STATS 531


Introduction to Flux

Dr Charles J Antonelli, LSA IT ARS
Mark Montague, LSA IT ARS
March, 2018

Preliminaries

Connecting to Flux

(Windows)

1. If you don't have PuTTY installed
Install the PuTTY/WinSCP installer for Windows:
Compute at the U: <http://its.umich.edu/computing/computers-software/compute>
Select + under Get Going
Get U-M PuTTY (UM_PuTTY_WinSCP.zip)
Execute the installer (UM_PuTTY_WinSCP.exe)
Accept all defaults
2. If you don't see this icon in your system tray
Install the Xming X Server for Windows
<https://sourceforge.net/projects/xming/files/Xming/6.9.0.31/> 
Get the Xming installer (Xming-6-9-0-31-setup.exe)
Execute the installer
Accept all defaults
<https://sourceforge.net/projects/xming/files/Xming-fonts/7.7.0.10/>
Get the Xming fonts installer (Xming-fonts-7-7-0-10-setup.exe)
Execute the installer
Accept all defaults

Connecting to Flux

(Windows)

1. Double-click "UM Internet Access Kit" icon on Desktop
Double-click PuTTY application within
2. In the Putty Application that appears:
Enter "flux-login.arc-ts.umich.edu" in the Host Name box
Under Connection | SSH | X11, ensure Enable X11 Forwarding is checked
Click "Open" at bottom
3. In the terminal window that appears:
Login with uniqname, Kerberos password, and Duo
4. This creates an ssh session on Flux

Connecting to Flux

(Linux & macOS)

1. Start a terminal window
2. In the terminal window that appears, type
`ssh -X username@flux-login.arc-ts.umich.edu`
3. Login to Flux with Kerberos password and Duo

Roadmap

- Preliminaries
- Introduction to Flux
- The command line
- Flux mechanics
- Parallel R and [pomp](#)

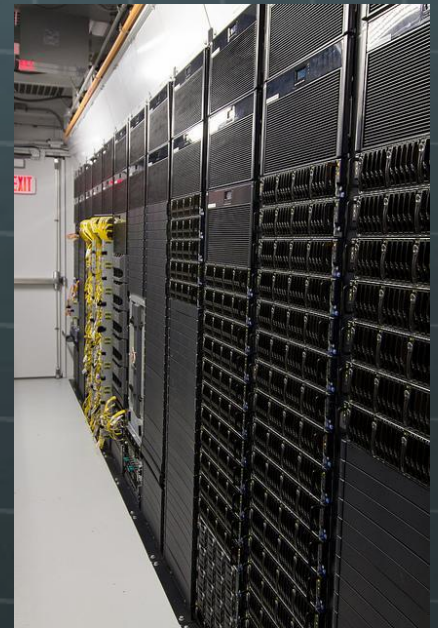
Introduction to Flux

Flux

Flux is a university-wide shared computational discovery / high-performance computing service.

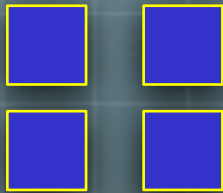
- Provided by Advanced Research Computing at U-M
- Procurement, licensing, billing by U-M ITS
- Interdisciplinary since 2010

<http://arc-ts.umich.edu/resources/compute-resources/>



The Flux cluster

Login nodes



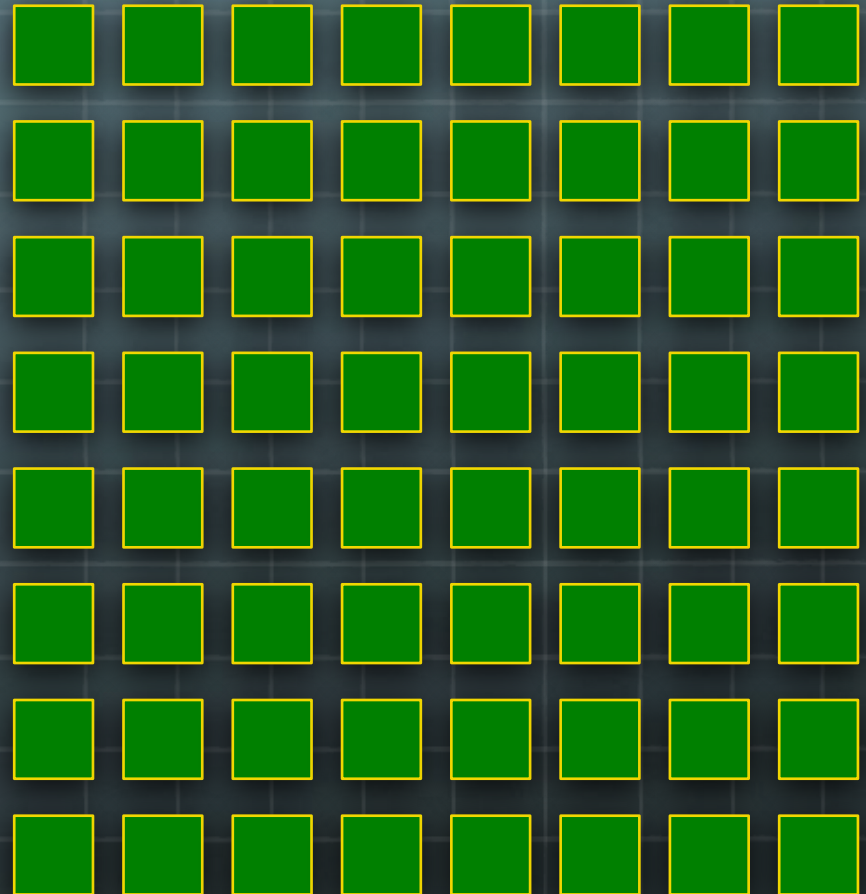
Data transfer
node



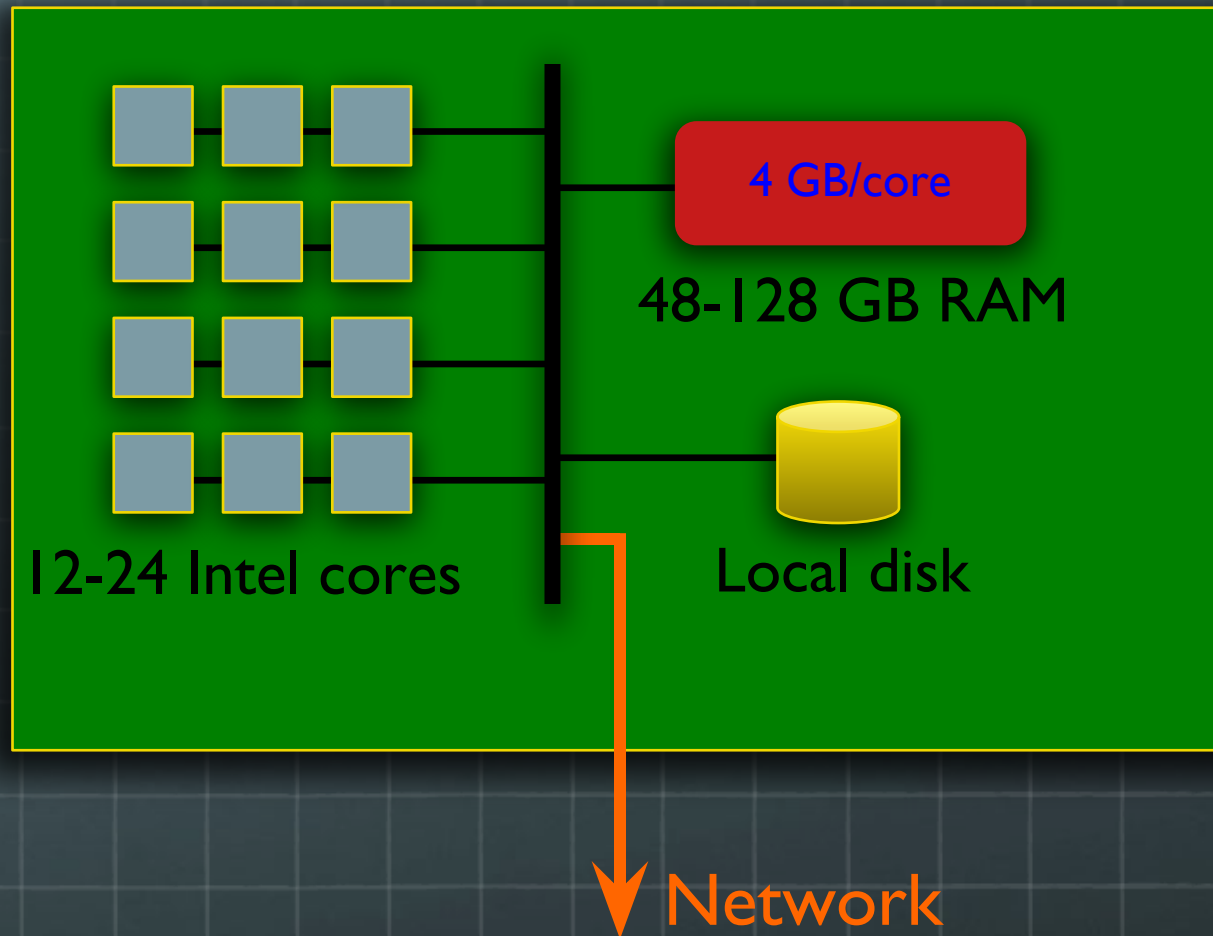
Storage



Compute nodes



A Standard Flux node



Using Flux

- Three basic requirements:
 - A Flux login account
<https://arc-ts.umich.edu/fluxform>
 - A Flux allocation
`stats531w18_flux`
 - A Duo app on your smartphone or tablet
<http://its.umich.edu/two-factor-authentication>
- Logging in to Flux
 - `ssh -X Login@flux-login.arc-ts.umich.edu`
 - PuTTY
 - Campus wired or MWireless
 - Otherwise use VPN, or
`ssh login.itd.umich.edu` first

The command line

Command Line Reference

William E Shotts, Jr.,

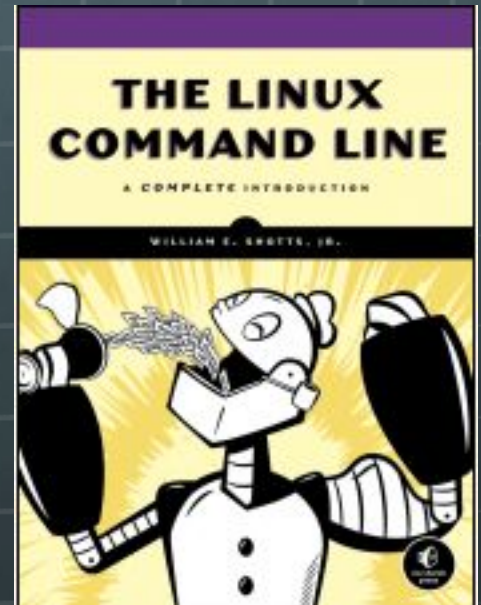
“The Linux Command Line: A Complete Introduction,”

No Starch Press, January 2012.

<http://linuxcommand.org/tlcl.php> .

Download Creative Commons Licensed
version at

[http://downloads.sourceforge.net/project/
linuxcommand/TLCL/13.07/TLCL-13.07.pdf](http://downloads.sourceforge.net/project/linuxcommand/TLCL/13.07/TLCL-13.07.pdf)



At the command prompt

- Basic input line editing commands
 - **Backspace** erases previous character
 - **Left** and **right arrow** move insertion point on the line
 - **Control-U** erases the line to the insertion point, so you can start over
 - **Enter** executes the line you typed
 - **Control-C** interrupts whatever command you started and returns you to the shell prompt (usually)
 - **Up** and **down arrow** will access your *command history*
 - Type “exit” without the quotes to exit the shell

Flux mechanics

Cluster batch workflow

- You create a batch script and submit it to PBS
- PBS schedules your job, and it enters the flux queue
- When its turn arrives, your job will execute the batch script
- Your script has access to all Flux applications and data
- When your script completes, anything it sent to standard output and error are saved in files stored in your submission directory
- You can ask that email be sent to you when your jobs starts, ends, or aborts
- You can check on the status of your job at any time, or delete it if it's not doing what you want
- A short time after your job completes, it disappears from PBS

Multi-threaded batch script

```
#PBS -N yourjobname  
#PBS -V  
#PBS -A youralloc_flux  
#PBS -q flux  
#PBS -l nodes=1:ppn=12,mem=47gb,walltime=00:05:00  
#PBS -m abe  
#PBS -j oe
```

```
#Your Code Goes Below:  
cd $PBS_O_WORKDIR  
R CMD BATCH --vanilla myscript.R myscript.out
```

Copying data

Using command line programs:

scp: copies files between hosts on a network over ssh

```
scp localfile username@flux-xfer.arc-ts.umich.edu:remotefile  
scp -r localdir username@flux-xfer.arc-ts.umich.edu:remotedir  
scp username@flux-login.arc-ts.umich.edu:remotefile localfile
```

Use "." as destination to copy to your Flux home directory:

```
scp localfile username@flux-xfer.arc-ts.umich.edu:.
```

... or to your Flux scratch directory:

```
scp localfile  
username@flux-xfer.arc-ts.umich.edu:/scratch/allocname/username
```

sftp: an interactive file transfer program over ssh (a secure ftp)

```
sftp username@flux-xfer.arc-ts.umich.edu
```

Using graphical (GUI) applications:

FileZilla (cross-platform): <http://filezilla-project.org/>

Cyberduck (Mac): <https://cyberduck.io/>

WinSCP (Windows): <http://www.itcs.umich.edu/bluedisc/>

Basic batch commands

- Once you have a script, submit it:

```
qsub scriptfile
```

```
$ qsub singlenode.pbs
6023521.nyx.engin.umich.edu
```

- You can check on the job status

```
qstat jobid
```

```
qstat -u user
```

```
$ qstat -u cja
nyx.engin.umich.edu:
```

Job ID	Username	Queue	Jobname	SessID	NDS	TSK	Req'd Memory	Req'd Time	Elap S	Elap Time
-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----
6023521.nyx.engi	cja	flux	hpc101i	--	1	1	--	00:05	Q	--

- To delete your job

```
qdel jobid
```


Modules

- The module command allows you to specify what Flux software you want to use

```
module list          -- Show loaded modules
module load name    -- Load module name for use
module avail         -- Show all available modules
module avail name   -- Show versions of module name
module key string   -- Search for string in module descrip
module spider string -- Search for string in all module doc
module unload name -- Unload module name
module use path     -- Add path to module search path
module              -- List all options
```

- Enter these commands at any time during your session
- You can load multiple modules in the same login session
- Software modules remain available throughout the session

Modules

- Once loaded, you can group a set of modules into a *module set*
`module save myset`
- A module set can be restored at any time
`module restore myset`
- List all module sets you've defined
`module savelist`
- Create a module set to be loaded each time you log in
`module save`

Lab

Task: Use the R multicore package

- Copy sample code to your login directory

```
cd
```

```
cp
```

```
/scratch/data/workshops/stats/stats-sample-code.tar.gz .
```

```
tar -zxvf stats-sample-code.tar.gz
```

```
cd ./stats-sample-code
```

- Examine `lab3.pbs` and `lab3.R`

Lab

Task: Use the R multicore package

- `module load R`
- Submit your job to Flux
`qsub lab3.pbs`
- Watch the progress of your job
`qstat -u username`
where *username* is your own username
- When complete, look at the job's output
`less lab3.out`

Parallel R and pomp

Random numbers

Serial pseudorandom sequence

- Usually based on linear recurrences modulo m

- Initialize the generator with some seed s

- Generate stream of pseudorandom numbers

Parallel pseudorandom sequence

- Performance, Reproducibility, Serializability

Install pomp

Task: Install pomp package

- `module load R/3.4.1`
- `R`
`install.packages("pomp", repos="https://cran.mtu.edu/")`
(answer y to "Would you like to use a personal library instead?" and "Would you like to create a personal library")

List of mirrors: <https://cran.r-project.org/mirrors.html>

Multicore example

```
rm(list=ls())  
library(doParallel)  
set.seed(2018,kind="L'Ecuyer")  
cores <- as.numeric(Sys.getenv('PBS_NP', unset='8'))  
cl <- makeCluster(cores)  
registerDoParallel(cl)  
trials <- 100
```

Multicore example

```
system.time(  
  r <- foreach(icount(trials),  
    .inorder=FALSE,  
    .options.multicore=list(set.seed=TRUE)  
  ) %dopar% {  
    library(pomp)  
    for (i in 1:100) {  
      pompExample("gompertz")  
      simulate(gompertz)  
    }  
  }  
)  
stopCluster(cl)
```

MPI example

```
system.time(  
  r <- foreach(icount(trials),  
    .inorder=FALSE,  
    .options.multicore=list(set.seed=TRUE)  
  ) %dopar% {  
    library(pomp)  
    for (i in 1:100) {  
      pompExample("gompertz")  
      simulate(gompertz)  
    }  
  }  
)  
stopCluster(cl)
```

Interactive jobs

- You can submit jobs *interactively*:

```
qsub -I -X -V -l nodes=1:ppn=2 -l walltime=15:00  
-A youralloc_flux -l qos=flux -q flux
```

- This queues a job as usual
 - Your terminal session will be blocked until the job runs
 - When your job runs, you'll get an interactive shell
 - When you exit the shell your job is deleted
- Interactive jobs allow you to
 - Develop and test on cluster node(s)
 - Execute GUI tools on a cluster node
 - Utilize a parallel debugger interactively

Lab

Task: Use an interactive PBS session

- `module load R`
- Start an interactive session
`qsub -I -V -l nodes=1:ppn=2 -l walltime=30:00 -A`
`stats531w18_flux -q flux`
- Run R in the interactive shell
`cd $PBS_O_WORKDIR`
`R`

Troubleshooting

- System-level
 - `freenodes` # aggregate node/core busy/free
 - `pbsnodes [-l]` # nodes, states, properties
with -l, list only nodes marked down
- Account-level
 - `mdiag -a acct` # cores & users for account *acct*
 - `showq [-r][-i][-b][-w acct=acct]` # running/idle/blocked jobs for *acct*
with -r|i|b show more info for that job state
 - `freealloc [--jobs] acct` # free resources in *acct*
with -jobs, shows resources in use
 - `idlenodes acct [property]` # shows available nodes for *acct* with *property*
- User-level
 - `mdiag -u uniq` # allocations for user *uniq*
 - `showq [-r][-i][-b][-w user=uniq]` # running/idle/blocked jobs for *uniq*
- Job-level
 - `qstat -f jobno` # full info for job *jobno*
 - `qstat -n jobno` # show nodes/cores where *jobno* running
 - `checkjob [-v] jobno` # show why *jobno* not running
 - `qpeek jobno` # peek at script output while *jobno* is running

Resources

- ARC User Guide <http://arc-ts.umich.edu/flux-user-guide/>
- ARC Flux pages <http://arc-ts.umich.edu/flux/>
- Software Catalog <http://arc-ts.umich.edu/software/>
- Quick Start Guide <http://arc-ts.umich.edu/flux/using-flux/flux-in-10-easy-steps/>
- Flux FAQs <http://arc-ts.umich.edu/flux/flux-faqs/>
- For assistance, send email to: hpc-support@umich.edu
 - Read by a team of people including unit support staff
 - Can help with Flux operational and usage questions
 - Programming support available